

Lab 6

Dynamic Memory Management

CSCI 4061
Introduction to Operating Systems

Oct 16, 2023

Overview

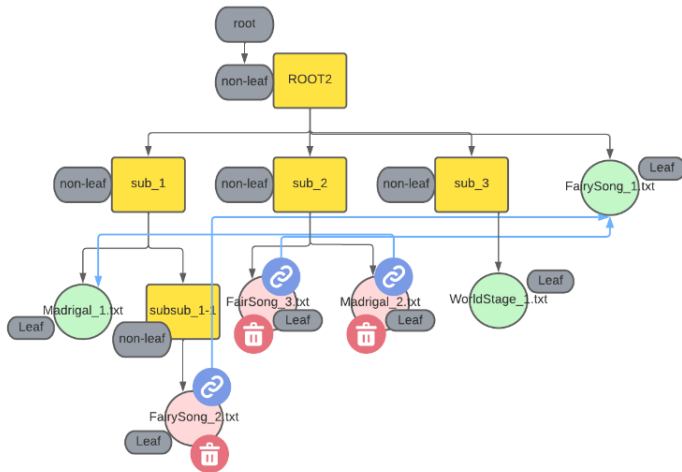
- PA2 Discussion
- Programming Exercise
- Memory allocation
- Valgrind

PA2 - Data deduplication

Project Goal

- ➊ Reduce the copies of data across machines and distributed file system to increase storage capacity
- ➋ Given a root folder with nested folders and files
- ➌ Copies of multiple files exists in the given folder hierarchy
- ➍ Find all copies and only maintain one
- ➎ The copies should be converted to symbolic link pointing to the maintained file

PA2 - Data deduplication



Template

- *root_process.c*
 - 1 Creates the first non-leaf process (for traversing root directory)
 - 2 The first non leaf process communicates all the aggregated file-hash from its children to the root process via a pipe
 - 3 The root process maintains unique files (one with smaller number) and deletes all duplicate files
 - 4 The paths of duplicate files points to the maintained file corresponding to them
 - 5 The root process will redirect stdout to given output file and report the symbolic link paths and its content (path to the original file)

Template

- *nonleaf_process.c*
 - 1 Traverse the given directory and performs two tasks
 - If the entry is a directory, create a child process and call `exec` on `nonleaf_process` with the directory entry
 - If the entry is a file, create a child process and call `exec` on `leaf_process` with the file entry
 - 2 The children should communicate the aggregated file-hash with the `nonleaf_process` via pipe
 - 3 The non-leaf process should aggregate all the file-hash from children and send it to its parent process via pipe

Template

- *leaf_process.c*
 - 1 leaf process is only called on files
 - 2 The leaf process will create hash from the given file using `hash_data_block` (`hash.h`)
 - 3 Send the file-hash to the parent process via pipe

PA2 - Data deduplication

Timeline

- ➊ Intermediate submission: Oct 20, 11:59 pm
- ➋ Final submission: Oct 27, 11:59 pm

Programming Exercise

Activity

- 1 You are given a program with multiple memory bugs
- 2 Resolve the bugs
- 3 Run valgrind on your code and copy results to a new file output.txt

We will cover memory allocation, memory leaks and using valgrind to identify memory leaks.

Memory Allocation

Static and dynamic allocation

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]){
    int arr_stat[10]; // static allocation of an array of size 10
    int *arr_dyn = (int *) malloc (sizeof(int) * 10); // dynamic allocation
                                                         // of an array of size 10

    for(int i = 0; i < 10; i++){
        arr_stat[i] = i;
        arr_dyn[i] = i;
    }

    printf("static size - %d\n", sizeof(arr_stat));
    printf("dynamic size - %d\n", sizeof(arr_dyn));

    free(arr_dyn); // frees allocated heap memory
    arr_dyn = NULL; // avoids double free
    return 0;
}
```

Memory Allocation

Terminal (p1.c)

```
$ make lab6_p1
./main
static size - 40
dynamic size - 8 # Reason for many errors in PA1
```

Memory Allocation

Terminal (p1.c)

```
$ make lab6_p1
./main
static size - 40
dynamic size - 8 # Reason for many errors in PA1
```

Solution

Malloc can be called only if we know size during runtime.
Keep track of this.

For strings you can use `strlen()` to get size of a char array

Memory allocation

Stack smashing

- 1 Stack grows downwards and Heap grows upwards in program memory layout
- 2 Storing more data than the memory allocated to a stack variable can cause buffer overflow access
- 3 You may end up corrupting adjacent data on the stack

Stack smash

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]){
    int arr[10];
    for(int i = 0; i < 11; i++){ // buffer overflow access
        arr[i] = i;
    }

    return 0;
}
```

Memory Allocation

Terminal (p2.c)

```
$ make lab6_p2
./main
*** stack smashing detected ***: terminated
make: *** [Makefile:7: lab6_p2] Aborted (core dumped)
```

Memory Allocation

Terminal (p2.c)

```
$ make lab6_p2
./main
*** stack smashing detected ***: terminated
make: *** [Makefile:7: lab6_p2] Aborted (core dumped)
```

Solution

Always check the memory size before storing the data to a buffer (bound check)

Memory allocation

Memory leaks / heap overflow

- 1 Heap memory allocations (malloc) should be explicitly freed in C
- 2 Unlike stack overflow, heap overflows are not reported as errors normally

Memory leak

```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[]){
    int *arr = (int *) malloc (sizeof(int) * 10);
    for(int i = 0; i < 11; i++){ // silent memory overflow
        arr[i] = i;
    }

    arr = NULL; // memory leak
    return 0;
}
```


Memory allocation

1. Memory leak identification

Use a memory mismanagement detector like Valgrind

Usage:

Compile you code with `-g` flag

```
valgrind --leak-check=full ./exe
```

Terminal (p3.c)

```
$ make lab6_p3_2
==363579== Invalid write of size 4
==363579==    at 0x1091AA: main (p3.c:7)
==363579== Address 0x4a99068 is 0 bytes after a block of size 40 alloc'd
==363579==    at 0x4848899: malloc (in /usr/libexec/valgrind/
                                                vgppreload_memcheck-amd64-linux.so)
==363579==    by 0x109185: main (p3.c:5)
==363579==
==363579==
==363579==
==363579== HEAP SUMMARY:
==363579==     in use at exit: 0 bytes in 0 blocks
==363579== total heap usage: 1 allocs, 1 frees, 40 bytes allocated
```

Memory allocation

2. Heap overflow error reporting

Use `-g` flag and `-fsanitize = address` when you compile code

Terminal (p3.c)

```
$ make lab6_p3_3
./main
=====
==371074==ERROR: AddressSanitizer: heap-buffer-overflow on address
          0x604000000038 at pc 0x559c2785e23f bp
          0x7ffd214eccf0 sp 0x7ffd214ecce0
WRITE of size 4 at 0x604000000038 thread T0
#0 0x559c2785e23e in main /home/csci4061/Lab6/code/sample/p3.c:7
#1 0x7fd58b01cd8f in __libc_start_call_main
    ./sysdeps/nptl/libc_start_call_main.h:58
#2 0x7fd58b01ce3f in __libc_start_main_impl ../csu/libc-start.c:392
#3 0x559c2785e104 in _start
    (/home/csci4061/Lab6/code/sample/main+0x1104)
```

Programming Exercise - Expected output

Terminal

```
$ make
==392853== Memcheck, a memory error detector
==392853== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==392853== Using Valgrind-3.18.1 and LibVEX; rerun with -h for copyright info
==392853== Command: ./main 3 3
==392853==
0 0 0
0 0 0
0 0 0
==392853==
==392853== HEAP SUMMARY:
==392853==      in use at exit: 0 bytes in 0 blocks
==392853==    total heap usage: 9 allocs, 9 frees, 1,144 bytes allocated
==392853==
==392853== All heap blocks were freed -- no leaks are possible
==392853==
==392853== For lists of detected and suppressed errors, rerun with: -s
==392853== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

Deliverables

Individual submission: Zip and submit to Gradescope by Oct 17, 11:59pm

- 1 bugs.c
- 2 output.txt